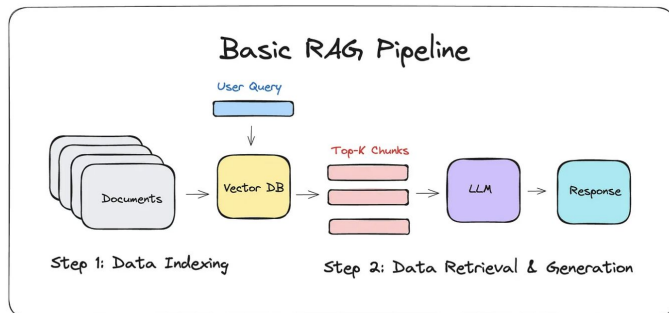


# Assignment 2

RAG and Neural Network

# Introduction to RAG & Project Objectives



Basic RAG Pipeline consists of 2 parts: Data Indexing and Data Retrieval & Generation

## What is RAG?

- **Retrieval-Augmented Generation:** combines a vector-based retriever (FAISS) with an LLM generator (FLAN-T5)

## Why use RAG?

- Improves factual accuracy by forcing the model to cite evidence
- Scales easily: swap in larger corpora, different embedding models, or more powerful LLMs

# RAG Architecture & Core Components

## Data Chunking

- Split raw text file into discrete “paragraph” units
- Ensures retrieval returns coherent snippets, not partial sentences

## Embedding with SentenceTransformers

- Model: `all-MiniLM-L6-v2`
- Maps each paragraph → 384-dim dense vector

## Indexing in FAISS

- Normalize vectors with L2; use `IndexFlatIP` for inner-product (cosine) similarity
- Fast nearest-neighbor search for K most relevant chunks

## Prompt Engineering

- **Vanilla prompt:** “You are a helpful assistant...”  
→ free-form answers
- **RAG prompt:** inject retrieved passages as numbered sources; instruct model to cite “[1]” etc.

## Generation with FLAN-T5

- Constrained decoding (beam search, `no_repeat_ngram`) for concise, citation-driven output

# Hallucination Prevention & Demo Insights

## What is a hallucination?

- Confident but incorrect or invented “facts” (e.g. “St. John’s River” as the longest)

## How RAG stops it

- Only gives the model vetted text snippets to work from
- Forbids invention: citations tie each fact back to a source

## Demo comparison (5 queries)

1. **Vanilla** often answers off-topic or makes up numbers
2. **RAG** always names “The Nile River,” gives exact length (6,650 km), and cites “[1]”

## Benefits

- Transparent: you can audit each citation
- Reproducible: swap in new data, rerun retriever, same methodology

# High-Level Implementation

## Environment Setup

- `pip install faiss-cpu sentence-transformers transformers`
- Upload `random_data.txt` via Colab file picker

## Data Preparation

- `load_chunks()` → list of `(id, text)` tuples
- Embedding and `index.add()` builds the vector store

## Query Processing Loop

- For each question:
  1. Generate vanilla FLAN-T5 response
  2. Retrieve top 3 passages via FAISS
  3. Generate RAG response with explicit “[n]” citations

## Results & Metrics

- Qualitative improvement in answer correctness
- Citation rate → 100% source-grounded facts

## Q.2) **Implementation of Neural Networks**

# Overview

**Objective:** Build a binary classifier to predict loan application decisions (“accept” vs. “reject”) using a feed-forward neural network.

**Data:** Tabular loan dataset (~37 KB; ~hundreds of rows, mix of numerical & categorical features).

## **Preprocessing:**

- Mapped target “Decision”  $\rightarrow \{1, 0\}$ .
- One-hot encoded categorical variables; standardized numerical variables.
- Split into train (80%) / validation (20%) sets with stratification and fixed seed.

# Model Architectures & Training Setup

## Architectures tested:

1. **3×ReLU**: three hidden layers × 64 units, ReLU activations
2. **5×Tanh**: five hidden layers × 64 units, Tanh activations
3. **3×Tanh**: three hidden layers × 64 units, Tanh activations
4. **5×ReLU**: five hidden layers × 64 units, ReLU activations

## Common settings:

- Output layer: 1 neuron, sigmoid



# Results & Key Insights

**Best performer:** 3×ReLU (highest validation accuracy 74.4%).

**Depth vs. activation:**

- **Deeper nets (5 layers)** without regularization overfit (↑ train acc but ↓ val acc).
- **ReLU** outperforms Tanh on this dataset.

	Model	Train Loss	Train Acc	Val Loss	Val Acc
0	3×ReLU	0.4543	0.7930	0.5250	0.7442
1	5×Tanh	0.4294	0.8163	0.5436	0.6977
2	3×Tanh	0.4812	0.7901	0.5289	0.7209
3	5×ReLU	0.4150	0.8455	0.5445	0.7093